



# Elasticsearch, You Know For Search! And More!

Medcl, 曾勇 (Zeng Yong)


Philips Kokoh Prasetyo  
Casey Vu  
Arinto Murdopo



**LARC**  
LIVING ANALYTICS  
RESEARCH CENTRE



**LARC**  
LIVING ANALYTICS  
RESEARCH CENTRE



Palanteer  
Systems

3.2M  
URLs  
Per  
Month

40M  
Tweets  
Per  
Month

14M  
Follow  
Links




Computer  
Cluster

# Outline

- Real time Search (Philips)
- Aggregation & Analytics (Casey)
- Lesson Learned @LARC (Arinto)

# Real-time Search



LUR (patrokan), prokaryotyczny instrument muzyczny o długości (do 2 m), wykonany z drewna, metalu, tworzywa sztucznego, w kształcie rurki z otworem w jednym końcu, z którego wydobywa się dźwięk. W Polsce LURy używano w czasie II wojny światowej w celu wywołania paniki w wrogu. W Polsce LURy używano w czasie II wojny światowej w celu wywołania paniki w wrogu.

LUR (patrokan), prokaryotyczny instrument muzyczny o długości (do 2 m), wykonany z drewna, metalu, tworzywa sztucznego, w kształcie rurki z otworem w jednym końcu, z którego wydobywa się dźwięk. W Polsce LURy używano w czasie II wojny światowej w celu wywołania paniki w wrogu.

LUR (patrokan), prokaryotyczny instrument muzyczny o długości (do 2 m), wykonany z drewna, metalu, tworzywa sztucznego, w kształcie rurki z otworem w jednym końcu, z którego wydobywa się dźwięk. W Polsce LURy używano w czasie II wojny światowej w celu wywołania paniki w wrogu.

LUR (patrokan), prokaryotyczny instrument muzyczny o długości (do 2 m), wykonany z drewna, metalu, tworzywa sztucznego, w kształcie rurki z otworem w jednym końcu, z którego wydobywa się dźwięk. W Polsce LURy używano w czasie II wojny światowej w celu wywołania paniki w wrogu.

LUR (patrokan), prokaryotyczny instrument muzyczny o długości (do 2 m), wykonany z drewna, metalu, tworzywa sztucznego, w kształcie rurki z otworem w jednym końcu, z którego wydobywa się dźwięk. W Polsce LURy używano w czasie II wojny światowej w celu wywołania paniki w wrogu.

LUR (patrokan), prokaryotyczny instrument muzyczny o długości (do 2 m), wykonany z drewna, metalu, tworzywa sztucznego, w kształcie rurki z otworem w jednym końcu, z którego wydobywa się dźwięk. W Polsce LURy używano w czasie II wojny światowej w celu wywołania paniki w wrogu.

LUR (patrokan), prokaryotyczny instrument muzyczny o długości (do 2 m), wykonany z drewna, metalu, tworzywa sztucznego, w kształcie rurki z otworem w jednym końcu, z którego wydobywa się dźwięk. W Polsce LURy używano w czasie II wojny światowej w celu wywołania paniki w wrogu.

LUR (patrokan), prokaryotyczny instrument muzyczny o długości (do 2 m), wykonany z drewna, metalu, tworzywa sztucznego, w kształcie rurki z otworem w jednym końcu, z którego wydobywa się dźwięk. W Polsce LURy używano w czasie II wojny światowej w celu wywołania paniki w wrogu.

LUR (patrokan), prokaryotyczny instrument muzyczny o długości (do 2 m), wykonany z drewna, metalu, tworzywa sztucznego, w kształcie rurki z otworem w jednym końcu, z którego wydobywa się dźwięk. W Polsce LURy używano w czasie II wojny światowej w celu wywołania paniki w wrogu.

LUR (patrokan), prokaryotyczny instrument muzyczny o długości (do 2 m), wykonany z drewna, metalu, tworzywa sztucznego, w kształcie rurki z otworem w jednym końcu, z którego wydobywa się dźwięk. W Polsce LURy używano w czasie II wojny światowej w celu wywołania paniki w wrogu.

LUR (patrokan), prokaryotyczny instrument muzyczny o długości (do 2 m), wykonany z drewna, metalu, tworzywa sztucznego, w kształcie rurki z otworem w jednym końcu, z którego wydobywa się dźwięk. W Polsce LURy używano w czasie II wojny światowej w celu wywołania paniki w wrogu.

# Real-time Search

- Indexing

- Mapping

- How a document, and the **fields** it contains are *stored* and *indexed*
      - Simple type: string, date, long, double, boolean, ip
      - Hierarchical: object, nested object
      - Specialized type: geo\_point, geo\_shape, completion

- Analysis

- How **full text** is *processed* to make it *searchable*
      - Standard Analyzer, Simple Analyzer, Language Analyzer, Snowball Analyzer

- Searching

- Query DSL

- Flexible and powerful query language used by Elasticsearch

# Mapping

```
{
  "id": 709260211836485600,
  "createdAt": "Mar 14, 2016 2:09:46 PM",
  "text": "@arinto Just few more days to share #elasticsearch at
#FOSSASIA 2016 https://t.co/PtZk14CNXl",
  "user": {
    "screenName": "philipskokoh",
    "name": "Philips Kokoh",
    ...
  },
  "hashtagEntities": [
    { "start": 36, "end": 50, "text": "elasticsearch" },
    { "start": 54, "end": 63, "text": "FOSSASIA" }
  ],
  "geoLocation": {
    "latitude": 1.2971576,
    "longitude": 103.8495769
  },
  ...
}
```

Long

Date

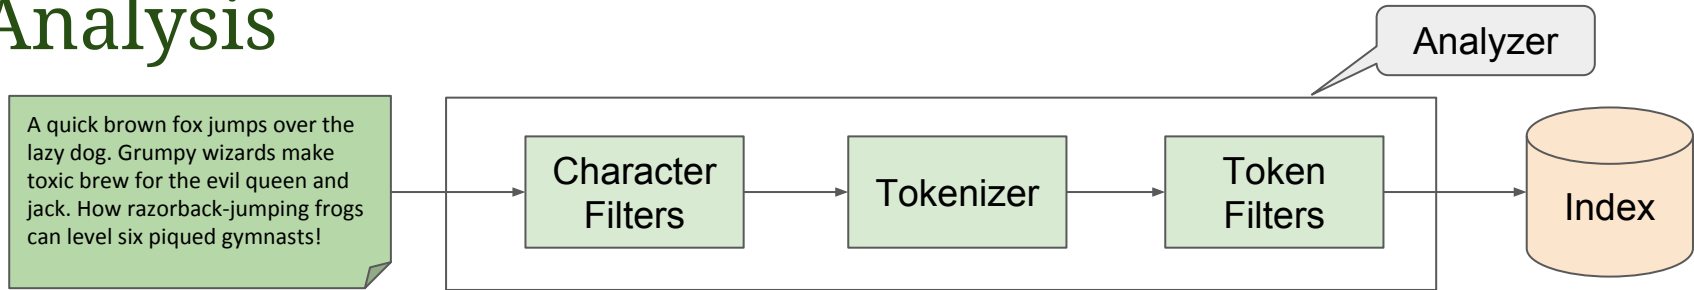
Analyzed String

Object

Nested Object

Geo\_point

# Analysis



## Built-in Analyzer

"Set the shape to semi-transparent by calling `set_trans(5)`"

- Standard Analyzer
  - [set, the, shape, to, semi, transparent, by, calling, set\_trans, 5]
- Simple Analyzer
  - [set, the, shape, to, semi, transparent, by, calling, set, trans]
- Whitespace Analyzer
  - [Set, the, shape, to, semi-transparent, by, calling, set\_trans(5)]
- Language Analyzer, e.g. english, french, spanish, arabic, ...
  - English analyzer: [set, shape, semi, transpar, call, set\_tran, 5]

# Searching

## Query DSL

- Based on JSON to define queries :)
- Behavior:
  - Query Context
    - Answering: “How well this document match this query clause?”
    - Return: **\_score** → **relevance score**
  - Filter Context
    - Answering: “Does this document match this query clause?”
    - Return: boolean **yes** or **no**



# Let's Search!!

Retrieve tweet containing "elasticsearch" **or** "fossasia"  
published before today by philipskokoh without geoLocation

```
{
  "query": {
    "bool": {
      "must": [
        {"match": { "text": "elasticsearch fossasia" }}
      ],
      "filter": [
        {"term": { "user.screenName": "philipskokoh" }},
        {"range": { "createdAt": {"lt": "now/d" }}}
      ],
      "must_not": [
        {"exists": { "field": "geoLocation" }}
      ]
    }
  }
}
```

# Let's Search!!

Retrieve tweet containing "elasticsearch" **or** "fossasia"  
published before today by philipskokoh without geoLocation

```
{
  "query": {
    "bool": {
      "must": [
        {"match": { "text": "elasticsearch fossasia" }}
      ],
      "filter": [
        {"term": { "user.screenName": "philipskokoh" }},
        {"range": { "createdAt": {"lt": "now/d" }}}
      ],
      "must_not": [
        {"exists": { "field": "geoLocation" }}
      ]
    }
  }
}
```

```
{
  "took": 22, "timed_out": false,
  "_shards": {
    "total": 42, "successful": 42, "failed": 0
  },
  "hits": {
    "total": 1, "max_score": 4.0619926,
    "hits": [
      {
        "_index": "plr_sg_tweet_201603",
        "_type": "tweet",
        "_id": "707403325390520320",
        "_score": 4.0619926,
        "_source": {
          ...
          "createdAt": "Mar 9, 2016 11:11:10 AM",
          "id": 707403325390520300,
          "text": "I will be giving a workshop at  
#FOSSASIA 2016 titled: Elasticsearch: You know, for  
search! and more! https://t.co/FRCQlQdHhH\nCome,  
join us!☐",
          "user": {
            ...
            "screenName": "philipskokoh",
            "lang": "en",
            "name": "Philips Kokoh",
          },
          "retweetCount": 2,
        }
      ]
    }
  }
}
```

# Geo Distance Range Query

Retrieve tweets containing fossasia published before today within 2km from Science Centre

```
{
  "query": {
    "bool": {
      "must": [
        {"match": { "text": "mrt" }}
      ],
      "filter": [
        {"range": { "createdAt": {"lt": "now/d" }}},
        {"geo_distance_range": {
          "gt": "0km", "lt": "1km",
          "geoLocation": {
            "lat": 1.332906,
            "lon": 103.736110
          }
        }}
      ]
    }
  }
}
```

# Geo Distance Range Query

Retrieve tweets containing fossasia published before today within 2km from Science Centre

```
{
  "query": {
    "bool": {
      "must": [
        {"match": { "text": "mrt" }}
      ],
      "filter": [
        {"range": { "createdAt": {"lt": "now/d" }}},
        {"geo_distance_range": {
          "gt": "0km", "lt": "1km",
          "geoLocation": {
            "lat": 1.332906,
            "lon": 103.736110
          }
        }}
      ]
    }
  }
}
```

```
{
  "took": 50, "timed_out": false,
  "_shards": {
    "total": 42, "successful": 42, "failed": 0
  },
  "hits": {
    "total": 974, "max_score": 3.6536937,
    "hits": [
      { ... },
      {
        "_index": "plr_sg_tweet_201602",
        "_type": "tweet",
        "_id": "700461563812192256",
        "_score": 3.646007,
        "_source": {
          ...
          "createdAt": "Feb 19, 2016 7:27:05 AM",
          "id": 700461563812192300,
          "text": "Mrt slow dao (@ Jurong East MRT
Interchange (NS1/EW24) - @smrt_singapore in
Singapore) https://t.co/K1av2dk7GI",
          "geoLocation": {
            "lat": 1.33378498,
            "lon": 103.74183655
          },
          "user": {
            "id": 252470398, ...
          }
        }
      },
      ...
    ]
  }
}
```

# Geo Bounding Box Query

Retrieve tweets containing singapore published inside Marina Bay area

```
{
  "query": {
    "bool": {
      "must": [
        {"match": { "text": "singapore" }}
      ],
      "filter": [
        {"geo_bounding_box": {
          "geoLocation": {
            "top_left": [103.852311, 1.289884],
            "bottom_right": [103.860465, 1.279158]
          }
        }}
      ]
    }
  }
}
```

Accept GeoJSON format!

```
{
  "took": 8, "timed_out": false,
  "_shards": {
    "total": 42, "successful": 42, "failed": 0
  },
  "hits": {
    "total": 5758, "max_score": 4.6547956,
    "hits": [
      {
        "_index": "plr_sg_tweet_201602",
        "_type": "tweet",
        "_id": "696276978584801280",
        "_score": 4.6547956,
        "_source": {
          ...
          "text": "In #Singapore",
          "geoLocation": {
            "lat": 1.28902587,
            "lon": 103.85594832
          },
          "user": { ... },
          ...
        }
      },
      ...
    ]
  }
}
```

# Nested Object

Retrieve tweets that starts with fossasia hashtag

```
{
  "query": {
    "nested": {
      "path": "hashtagEntities",
      "query": {
        "bool": {
          "must": [
            {"match": { "hashtagEntities.text": "fossasia"}}
          ],
          "filter": [
            {"range": { "hashtagEntities.start": { "lt": 1
            }}}
          ]
        }
      }
    }
  }
}
```

# Nested Object

Retrieve tweets that starts with fossasia hashtag

```
{
  "query": {
    "nested": {
      "path": "hashtagEntities",
      "query": {
        "bool": {
          "must": [
            {"match": { "hashtagEntities.text": "fossasia"}}
          ],
          "filter": [
            {"range": { "hashtagEntities.start": { "lt": 1
            }}}
          ]
        }
      }
    }
  }
}
```

```
{
  "took": 6, "timed_out": false,
  "_shards": {
    "total": 42, "successful": 42, "failed": 0
  },
  "hits": {
    "total": 3, "max_score": 16.199848,
    "hits": [
      { "_score": 16.199848, ...
        "_source": {
          ...
          "hashtagEntities": [
            { "start": 0, "end": 9,
              "text": "FOSSASIA" }
          ],
          "text": "#FOSSASIA #GoogleCodeIn #GCI
speeding up. 150+ students currently working on
tasks! Great you are joining @hpdang
@mariobehling @mohitkanwal"
        }
      }, { "_score": 15.867135, ...
        "_source": {
          ...
          "hashtagEntities": [
            { "start": 0, "end": 9,
              "text": "FOSSASIA" },
          ],
          "text": "#FOSSASIA 2016 is keen to
get more students to attend. Learn coding n
tech. Happy to share more details https://t.
co/1bRmvZVrOP #edsg"
        }
      }, { ... }
    ]
  }
}
```



# Aggregation & Analytics



# Aggregation and Analytics

Types of aggregations (that we often use):

- Terms Aggregation:
  - Bucketing documents based on numeric/textual content
- Date Histogram Aggregation:
  - Bucketing documents based on date/time value
- Geo Distance Aggregation
  - Bucketing documents based on distance from an origin location

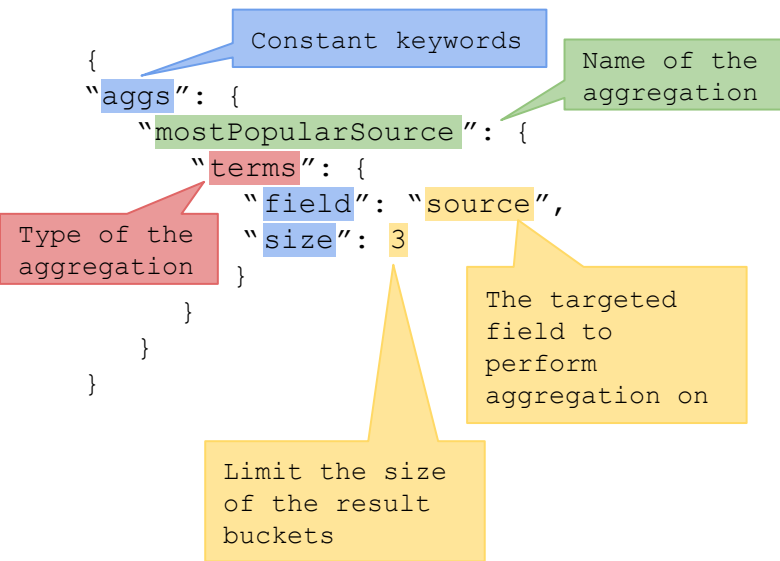
Combined Aggregations

Combined Aggregations & Queries

# Terms Aggregation

*“What are the popular platforms Twitter users use?”*

# Terms Aggregation



# Terms Aggregation

```
{
  "aggs": {
    "mostPopularSource": {
      "terms": {
        "field": "source",
        "size": 3
      }
    }
  }
}
```

```
{ ...
  "aggregations": {
    "mostPopularSource": {
      "doc_count_error_upper_bound": 87696,
      "sum_other_doc_count": 12907898,
      "buckets": [
        {
          "key": "Twitter for iPhone",
          "Doc_count": 27928770
        },
        {
          "key": "Twitter for Android",
          "Doc_count": 21327691
        },
        {
          "key": "Twitter Web Client",
          "Doc_count": 6243422
        }
      ]
    }
  }
}
```

Doc counts are approximate (-> upper bound on doc\_count error for each term)

The sum of doc counts for buckets not in the response



Term: the bucket's keyword

The doc counts for this bucket

# Date Histogram Aggregation

*“Number of tweets collected each month?”*

# Date Histogram Aggregation

```
{  
  "aggs": {  
    "numberOfTweetsByMonth": {  
      "date_histogram": {  
        "field": "createdAt",  
        "interval": "month"  
      }  
    }  
  }  
}
```

Type of the  
aggregation

Define the  
interval to  
"bucket" the  
count

The targeted  
field to  
perform  
aggregation on

# Date Histogram Aggregation

```
{
  "aggs": {
    "numberOfTweetsByMonth": {
      "date_histogram": {
        "field": "createdAt",
        "interval": "month"
      }
    }
  }
}
```



```
{ ...
  "aggregations": {
    "numberOfTweetsByMonth": {
      "buckets": [
        {
          "key_as_string":
            "Jan 1, 2016 12:0:0 AM",
          "key": 1451606400000
          "doc_count": 28067435
        },
        {
          "key_as_string":
            "Feb 1, 2016 12:0:0 AM",
          "key": 1454284800000
          "doc_count": 25912385
        },
        {
          "key_as_string":
            "Mar 1, 2016 12:0:0 AM",
          "key": 1456790400000
          "doc_count": 14427961
        },
        ..., ...}}}
}
```

We have 28 millions tweets that was tweeted (createdAt) in Jan 2016

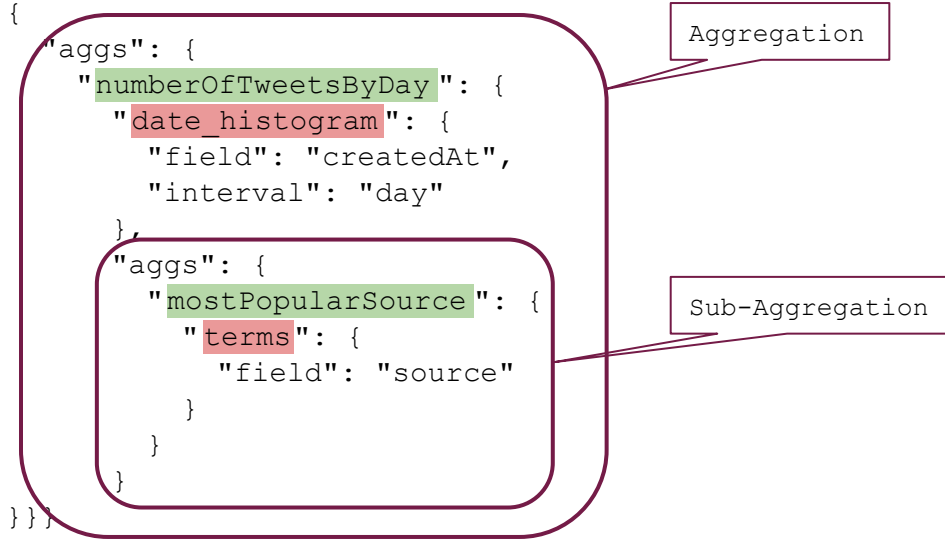
# Terms + Date Histogram Aggregation Combined

*“What is the platform Twitter users use?”*


*“On each day”*



# Terms + Date Histogram Aggregation Combined



# Terms + Date Histogram Aggregation Combined



# Geo Distance Aggregations

```
{
  "aggs": {
    "numberOfTweetsByRadius": {
      "geo_distance": {
        "field": "geoLocation",
        "origin": "1.3,103.8",
        "unit": "km",
        "ranges": [
          { "to": 1 },
          { "from": 1, "to": 3 },
          { "from": 3 }
        ]
      }
    }
  }
}
```

Type of the aggregation

The targeted field

Distances are computed from this origin

The radius' unit

Define the buckets

# Geo Distance Aggregations

```
{
  "aggs": {
    "numberOfTweetsByRadius": {
      "geo_distance": {
        "field": "geoLocation",
        "origin": "1.3,103.8",
        "unit": "km",
        "ranges": [
          { "to": 1 },
          { "from": 1, "to": 3 },
          { "from": 3 }
        ]
      }
    }
  }
}
```



```
{ ...
  "aggregations": {
    "numberOfTweetsByRadius": {
      "buckets": [
        {
          "key": "*-1.0",
          "from": 0, "from_as_string": "0.0",
          "to": 1, "to_as_string": "1.0"
          "doc_count": 1578
        },
        {
          "key": "1.0-3.0",
          "from": 1, "from_as_string": "1.0",
          "to": 3, "to_as_string": "3.0",
          "doc_count": 17880
        },
        {
          "key": "3.0-*",
          "from": 3, "from_as_string": "3.0"
          "doc_count": 779613
        }
      ]
    }
  }
}
```

# Geo Distance Aggregations



# And many other types of aggregations

Refer to:




<https://www.elastic.co/guide/en/elasticsearch/reference/current/search-aggregations.html>

- Aggregations
  - Metrics Aggregations
    - Avg Aggregation
    - Cardinality Aggregation
    - Extended Stats Aggregation
    - Geo Bounds Aggregation
    - Geo Centroid Aggregation
    - Max Aggregation
    - Min Aggregation
    - Percentiles Aggregation
    - Percentile Ranks Aggregation
    - Scripted Metric Aggregation
    - Stats Aggregation
    - Sum Aggregation
    - Top hits Aggregation
    - Value Count Aggregation
  - + Bucket Aggregations
  - + Pipeline Aggregations
    - Caching heavy aggregations
    - Returning only aggregation results
    - Aggregation Metadata


- Aggregations
  - + Metrics Aggregations
  - Bucket Aggregations
    - Children Aggregation
    - Date Histogram Aggregation
    - Date Range Aggregation
    - Filter Aggregation
    - Filters Aggregation
    - Geo Distance Aggregation
    - GeoHash grid Aggregation
    - Global Aggregation
    - Histogram Aggregation
    - IPv4 Range Aggregation
    - Missing Aggregation
    - Nested Aggregation
    - Range Aggregation
    - Reverse nested Aggregation
    - Sampler Aggregation
    - Significant Terms Aggregation
    - Terms Aggregation
  - + Pipeline Aggregations

- Aggregations
  - + Metrics Aggregations
  - + Bucket Aggregations
  - Pipeline Aggregations
    - Avg Bucket Aggregation
    - Derivative Aggregation
    - Max Bucket Aggregation
    - Min Bucket Aggregation
    - Sum Bucket Aggregation
    - Stats Bucket Aggregation
    - Extended Stats Bucket Aggregation
    - Percentiles Bucket Aggregation
    - Moving Average Aggregation
    - Cumulative Sum Aggregation
    - Bucket Script Aggregation
    - Bucket Selector Aggregation
    - Serial Differencing Aggregation
  - Caching heavy aggregations
  - Returning only aggregation results
  - Aggregation Metadata

in Beta



# Lesson Learned






# ES @ LARC: Lesson Learned

0.19.x -> 0.90.10 -> 1.7.0 -> 2.0.0



# Read

Elasticsearch the Definitive Guide ([online](#))



Foreword

+ Preface

- Getting Started

- + You Know, for Search...
- + Life Inside a Cluster
- + Data In, Data Out
- + Distributed Document Store
- + Searching—The Basic Tools
- + Mapping and Analysis
- + Full-Body Search
- + Sorting and Relevance
- + Distributed Search Execution
- + Index Management
- + Inside a Shard

Start with Getting Started section!

+ Search in Depth

+ Dealing with Human Language

+ Aggregations

+ Geolocation

+ Modeling Your Data

+ Administration, Monitoring, and Deployment

# Define the correct mapping

Elasticsearch dynamic mapping, OK for production?

```
{
  "id": 709260211836485600,
  "createdAt": "Mar 14, 2016 2:09:46 PM",
  "text": "@arinto Just few more days to
share #elasticsearch at
#FOSSASIA 2016 https://t.co/PtZk14CNXl",
  "geoLocation": {
    "latitude": 1.2971576,
    "longitude": 103.8495769
  },
  "favorites": [{
    "screenName": "arinto",
    "origin": "Indonesia"},
    {
    "screenName": "casey",
    "origin": "Vietnam"}]
}
```

string instead of date

double instead of geo\_point

No relation between fields!  
Searching for  
(screenName == arinto && origin == vietnam)  
will return both data

# Define the correct mapping

```
{
  "id": 709260211836485600,
  "createdAt": "Mar 14, 2016 2:09:46 PM",
  "text": "@arinto Just few more days to
    share #elasticsearch at
    #FOSSASIA 2016 https://t.co/PtZk14CNX1",
  "geoLocation": {
    "latitude": 1.2971576,
    "longitude": 103.8495769
  },
  "favorites": [{
    "screenName": "arinto",
    "origin": "Indonesia"},
    {
    "screenName": "casey",
    "origin": "Vietnam"}]
}
```

```
{
  //rest of the mapping
  "id": {
    "type": "long" },
  "createdAt": {
    "format": "MMM d, y h:m:s a",
    "type": "date" },
  "text": {
    "type": "string" },
  "geoLocation": {
    "type": "geo_point" },
  "favorites": {
    "type": "nested"
    "properties": {...}}
  //..rest of the mapping
}
```

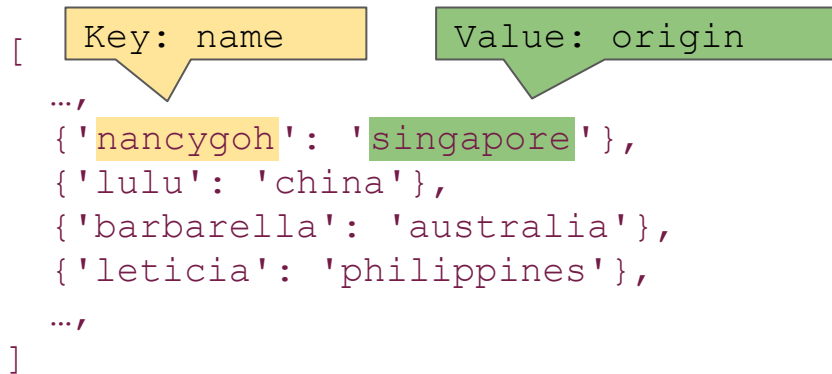
Elasticsearch dynamic mapping, not OK for production.

**Define the correct mapping!** Check the docs to learn more about mapping

# KV store in Elasticsearch

## Key-value store in Elasticsearch

- Field name as the key
- 10 or 100 keys are okay..
- What if you have million of keys?
- Does it scale?



The diagram illustrates a key-value store structure. It shows a list of objects, each containing a key-value pair. The key is highlighted in yellow and labeled 'Key: name', and the value is highlighted in green and labeled 'Value: origin'. The objects are enclosed in brackets and separated by commas.

```
[  
  ...,  
  {'nancygoh': 'singapore'},  
  {'lulu': 'china'},  
  {'barbarella': 'australia'},  
  {'leticia': 'philippines'},  
  ...,  
]
```

# KV store - Mapping Explosion!

- Dynamically add new fields in a mapping is an **expensive** operation
  - Lock the index, add new fields, and propagate index structure changes
- Halt the cluster!



# KV store - Solution

## Correct mapping:

```
{
  "mydata": {
    "mappings": {
      "kv": {
        "dynamic": "strict",
        "properties": {
          "key": {
            "type": "string",
            "index": "not_analyzed" },
          "value": {
            "type": "nested",
            "properties": {.....} //detail hidden
          }
        }
      }
    }
  }
}
```

## Sample indexed data:

```
[{
  "key": "nancygoh",
  "value": { "origin": "singapore " }
}, {
  "key": "lulu",
  "value": { "origin": "china" }
}, {
  "key": "barbarella",
  "value": { "origin": "australia" }
}, {
  "key": "leticiabongnino",
  "value": { "origin": "philippines" }
}]
```

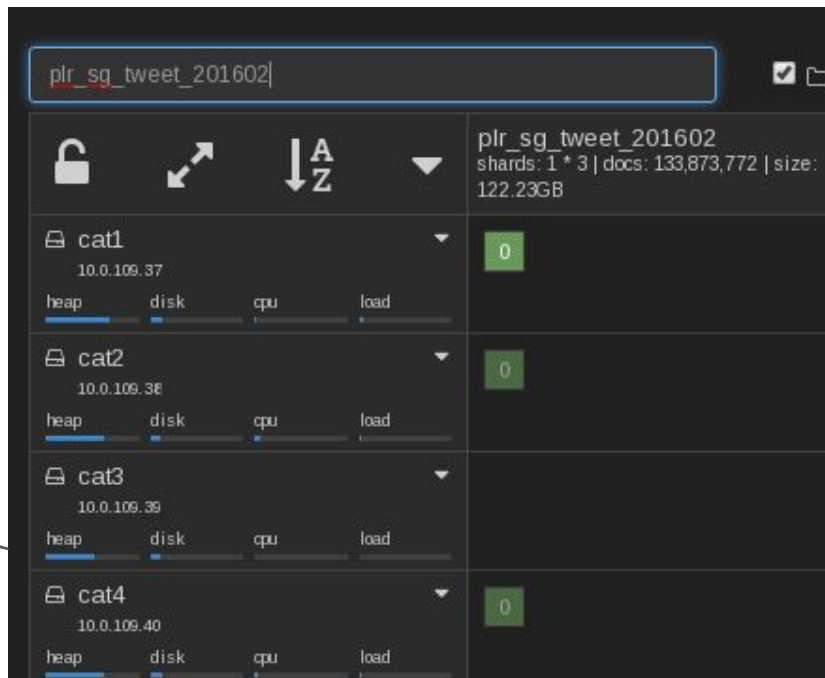
**Keep the number of fields under control!**



# Shards

- No magic number
  - However.. you must determine when you create the index
  - Estimate data growth
  - Prepare for reindex

kopf plugin



# Index template & rolling index

Rolling (time-based)  
index

CREATE INDEX

Index template

Pattern: plr\_sg\_tweet\_\*

template name

plr\_sg\_tweet

body

```
1 {
2   "order": 1,
3   "template": "plr_sg_tweet_*",
4   "settings": {
5     "index": {
6       "number_of_shards": "1",
7       "number_of_replicas": "1"
8     }
9   },
10  "mappings": {
11    "tweet": {
12      "dynamic": "strict",
13      "properties": {
14        "sentiment": {
15          "type": "long"
16        }
17      }
18    }
19  }
20 }
```

Reconfigure  
shards &  
replicas

Define the  
correct mapping

plr\_sg\_tweet\_201602



plr\_sg\_tweet\_201602  
shards: 1 \* 3 | docs: 133,873,772 | size:  
122.23GB

cat1

10.0.109.37

heap

disk

cpu

load

0

cat2

10.0.109.38

heap

disk

cpu

load

0

cat3

10.0.109.39

heap

disk

cpu

load

0

cat4

10.0.109.40

heap

disk

cpu

load

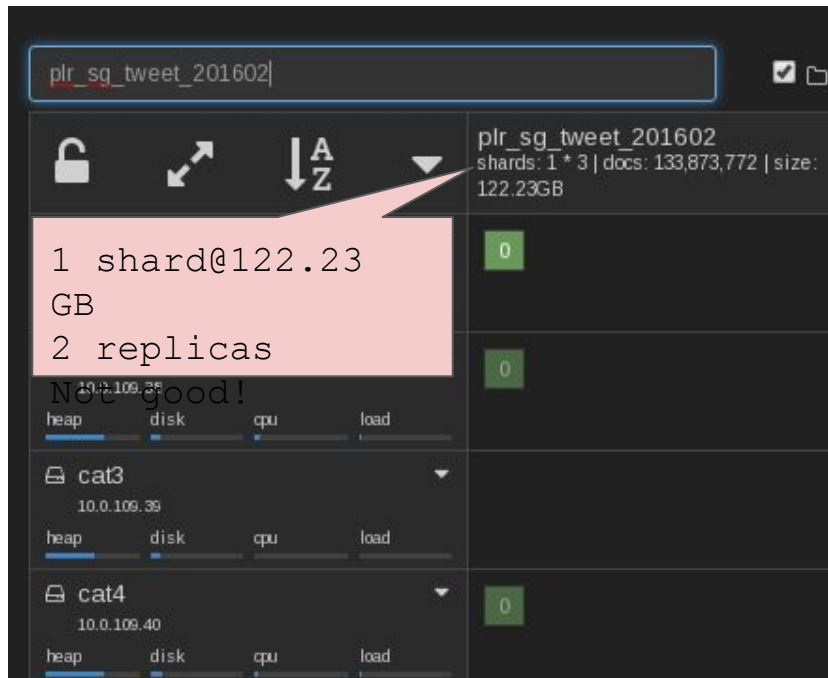
0

# Problematic shard

```
CREATE INDEX TEMPLATE

template name
plr_sg_tweet

body
1 {
2   "order": 1,
3   "template": "plr_sg_tweet_*",
4   "settings": {
5     "index": {
6       "number_of_shards": "1",
7       "number_of_replicas": "1"
8     }
9   },
10  "mappings": {
11    "tweet": {
12      "dynamic": "strict",
13      "properties": {
14        "sentiment": {
15          "type": "long"
16        }
17      }
18    }
19  }
20 }
```



# Modify index template

```
CREATE INDEX
Index template
Pattern:
template name: plr_sg_tweet_*
plr_sg_tweet

body
1 {
2   "order": 1,
3   "template": "plr_sg_tweet_*",
4   "settings": {
5     "index": {
6       "number_of_shards": "14",
7       "number_of_replicas": "1"
8     }
9   },
10  "mappings": {
11    "tweet": {
12      "dynamic": "strict",
13      "properties": {
14        "sentiment": {
15          "type": "long"
16        }
17      }
18    }
19  }
20 }
```


Reconfigure  
shards &  
replicas

201603  
14 shards  
1 replica

plr_sg_tweet_201603				plr_sg_tweet_201603			
				shards: 14 * 2   docs: 89,500,925   size: 96.33GB			
cat1 10.0.109.37				0 4 6 7			
heap disk cpu load							
cat2 10.0.109.38				1 2 5 8			
heap disk cpu load							
cat3 10.0.109.39				6 9 10 11			
heap disk cpu load							
cat4 10.0.109.40				2 3 9 11			
heap disk cpu load							
cat5							

# Indexing performance

## Tweet bulk-loading indexing performance



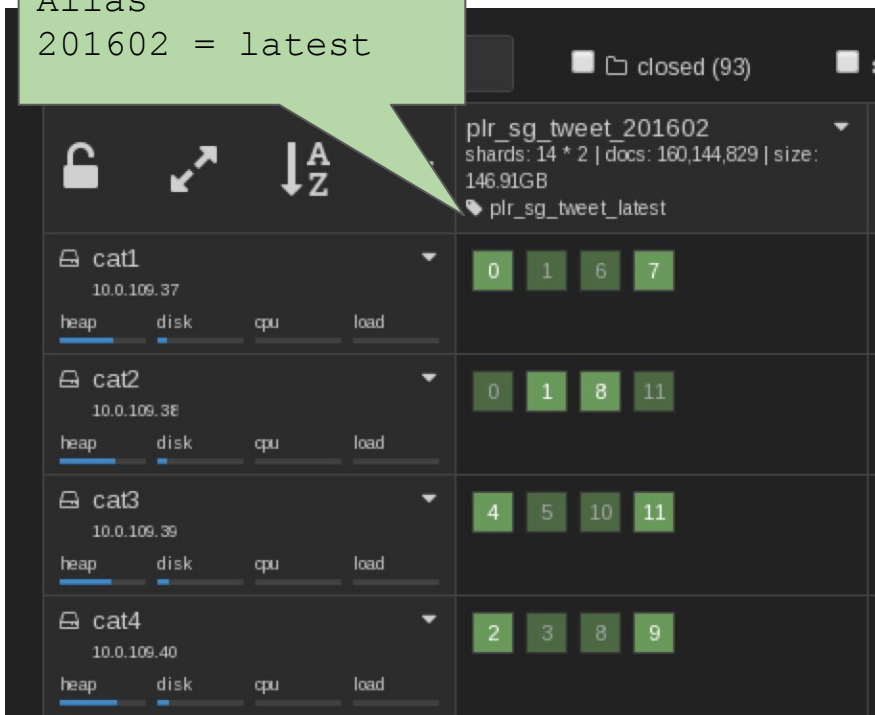
# Alias

No change in application code:

```
1 POST plr_sg_tweet_latest/tweet/_search
2 {
3   "query": {"match_all": {}},
4   "size": 3
5 }
```

```
1 {
2   "took": 37,
3   "timed_out": false,
4   "_shards": {
5     "total": 14,
6     "successful": 14,
7     "failed": 0
8   },
9   "hits": {
10    "total": 25912385,
11    "max_score": 1,
12    "hits": [
13      {
14        "_index": "plr_sg_tweet_201602",
15        "_type": "tweet",
16        "_id": "696897598221758464",
17        "_score": 1,
18        "_source": {
19          "inReplyToUserId": 371748246,
```

Alias  
201602 = latest



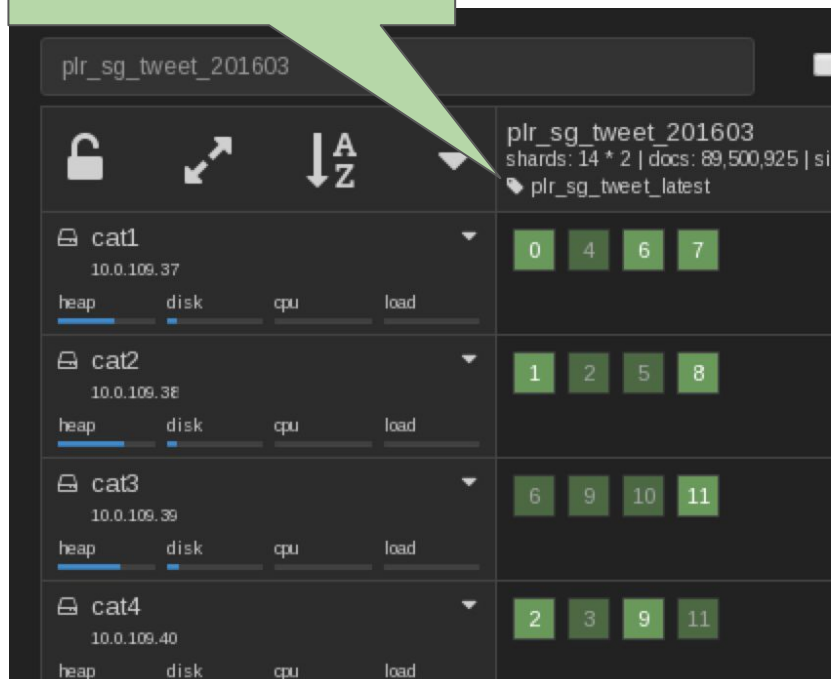
# Alias

## No change in application code:

```
1 POST plr_sg_tweet_latest/tweet/_search
2 {
3   "query": {"match_all": {}},
4   "size": 3
5 }
```

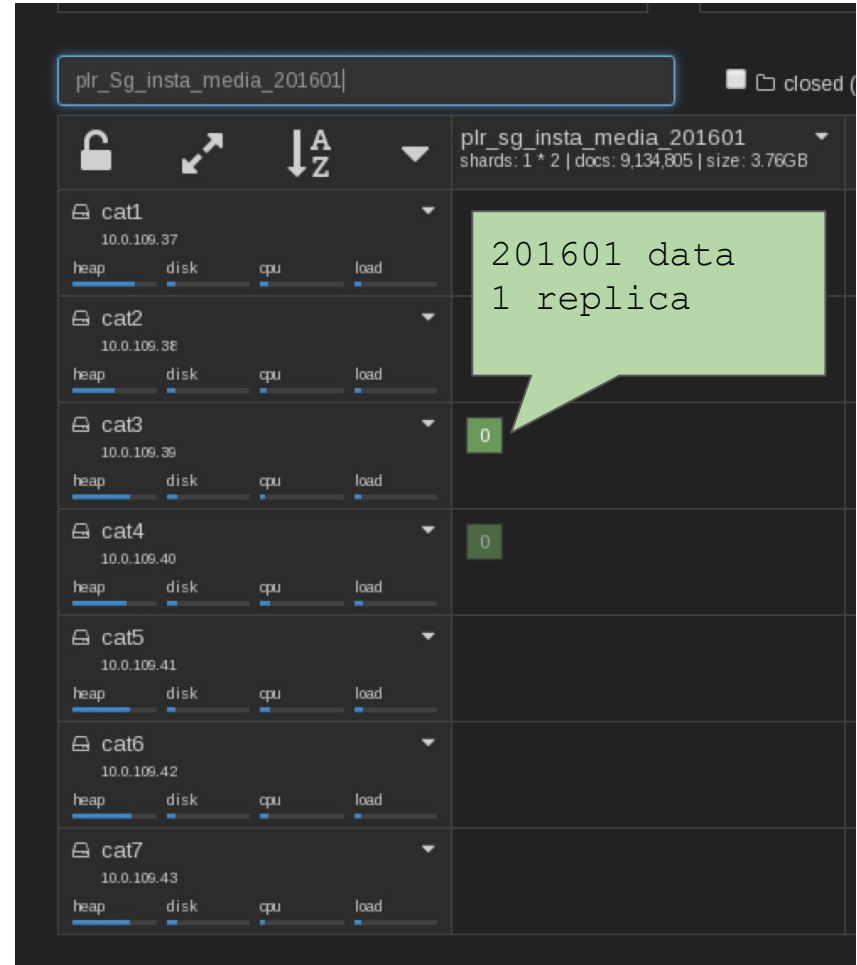
```
1 {
2   "took": 27,
3   "timed_out": false,
4   "_shards": {
5     "total": 14,
6     "successful": 14,
7     "failed": 0
8   },
9   "hits": {
10    "total": 14427961,
11    "max_score": 1,
12    "hits": [
13      {
14        "_index": "plr_sg_tweet_201603",
15        "_type": "tweet",
16        "_id": "706690831164387328",
17        "_score": 1,
18        "source": {
```

```
Alias
201603 = latest
```



# Replicas

- You can change it on the fly
- Start with 1 for fault tolerance
- What happen when the read request rate is very high?





# Replicas

- Increase accordingly to
  - balance load
  - increasing the availability
  - scale up read requests



# Lesson Learned

- Read the book
- Define correct mapping
- Index template & rolling indices
- Use `alias`
- Scale up using replicas

# Q&A



## Elastic Training in Singapore:

- Core Elasticsearch: Operations 25 April 2016
- Core Elasticsearch: Developer 26-27 April 2016

[training.elastic.co](http://training.elastic.co)



# We're Hiring!!!

[bit.ly/larchiring](http://bit.ly/larchiring)