




Unlocking New Big Data Insights with Hadoop & MySQL

Ricky Setyawan
MySQL Principal Consultant - ASEAN

An Avalanche of Data

**3B to 50B
Devices**

Source: Ericsson




Big Data What It Is, What it Means

Volume

Variety

Velocity



Create Value


What's Changed?

- Enablers
 - Digitization – nearly everything has a digital heartbeat
 - Ability to store much larger data volumes (distributed file systems)
 - Ability to process much larger data volumes (parallel processing)
- Why is this different from BI/DW?
 - Business formulated questions to ask upfront
 - Drove what was data collected, data model, query design

➤ Big Data Enables what-if analysis, real-time discovery

Big Data Adoption

- Web Recommendations
- Sentiment Analysis
- Marketing Campaign Analysis
- Customer Churn Modeling
- Fraud Detection
- Research and Development
- Risk Modeling
- Machine Learning




Leading Use-Case, On-Line Retail




Why Hadoop?

- Scales to thousands of nodes, PB of structured and unstructured data
 - Combines data from multiple sources, schema-less
 - Run queries against all of the data
- Runs on commodity servers, handle storage and processing
- Data replicated, self-healing
- Initially just batch (Map/Reduce) processing
 - Extending with interactive querying, via Apache Drill, Cloudera Impala, Stinger etc.




Big Data Lifecycle

Better Decisions Using Big Data



Big Data Lifecycle

Better Decisions Using Big Data




MySQL Database
MySQL Cluster
JSON Support
NoSQL Interfaces
MySQL Fabric

MySQL 5.7 Sysbench Benchmark: SQL Point Selects

3x Faster than MySQL 5.6

1,600,000 QPS

MySQL 5.7: Sysbench OLTP Read Only (SQL Point Selects)



Intel(R) Xeon(R) CPU E7-8890 v3
4 sockets x 18 cores-HT (144 CPU threads)
2.5 Ghz, 512GB RAM
Linux kernel 3.16

Hybrid Database: Rock Solid Reliability + Flexibility

Modern Applications

Require **agile** development and operations with **robust data protection** and **security**

Traditional RDBMS
Proven, transactional, secure
Complex JOINS and queries
Extensive **operational** tools

NoSQL Solutions
Flexible. Easy-to-use.
Schema-less document storage




MySQL 5.7
JSON Support

Hybrid Database

No trade-offs, best of both worlds. ACID properties & reliability of RDMS + flexible document management

MySQL NoSQL Interfaces: Fast, Flexible, Safe

Blazing Fast
Key/Value Queries



Fully Transactional/
ACID

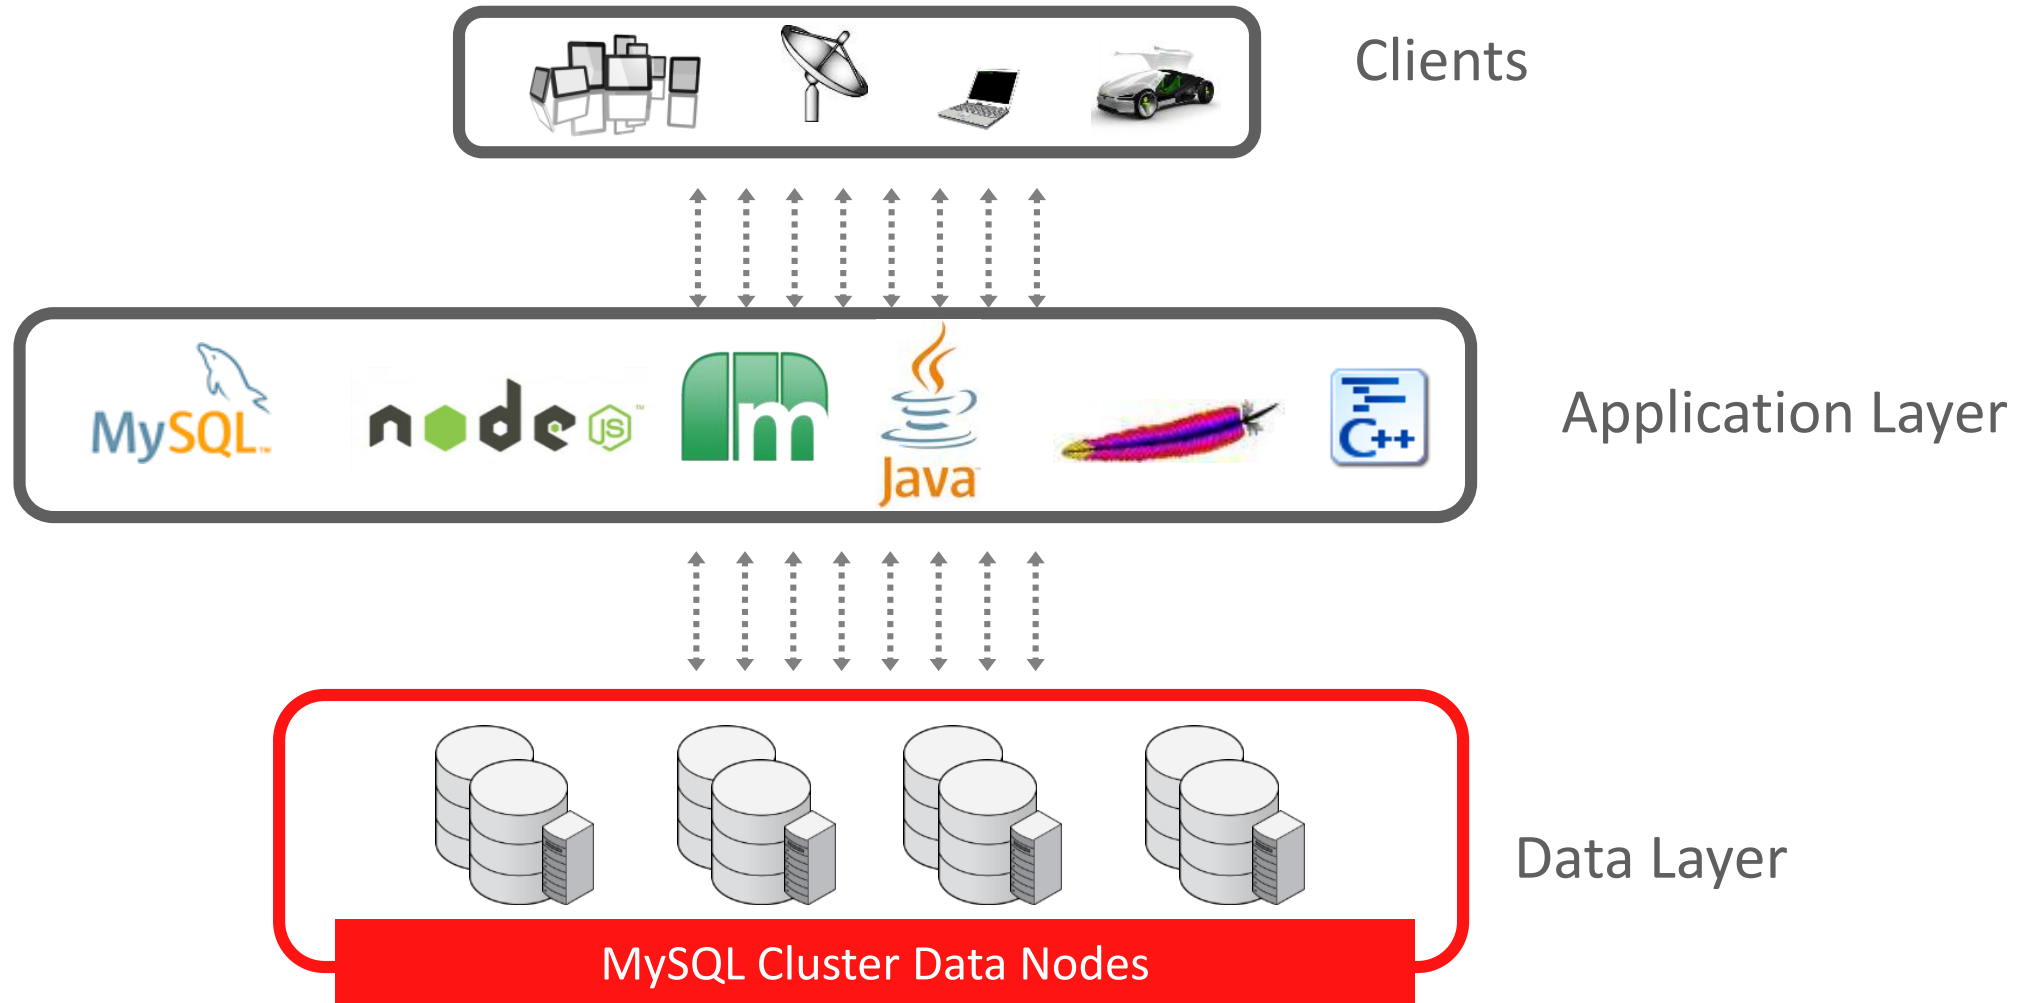


NoSQL And SQL
Across the same
data Set



Combined with Schema Flexibility: Online DDL


NoSQL Interfaces to MySQL Cluster



MySQL Cluster 7.4 NoSQL Performance

200 Million NoSQL Reads/Second

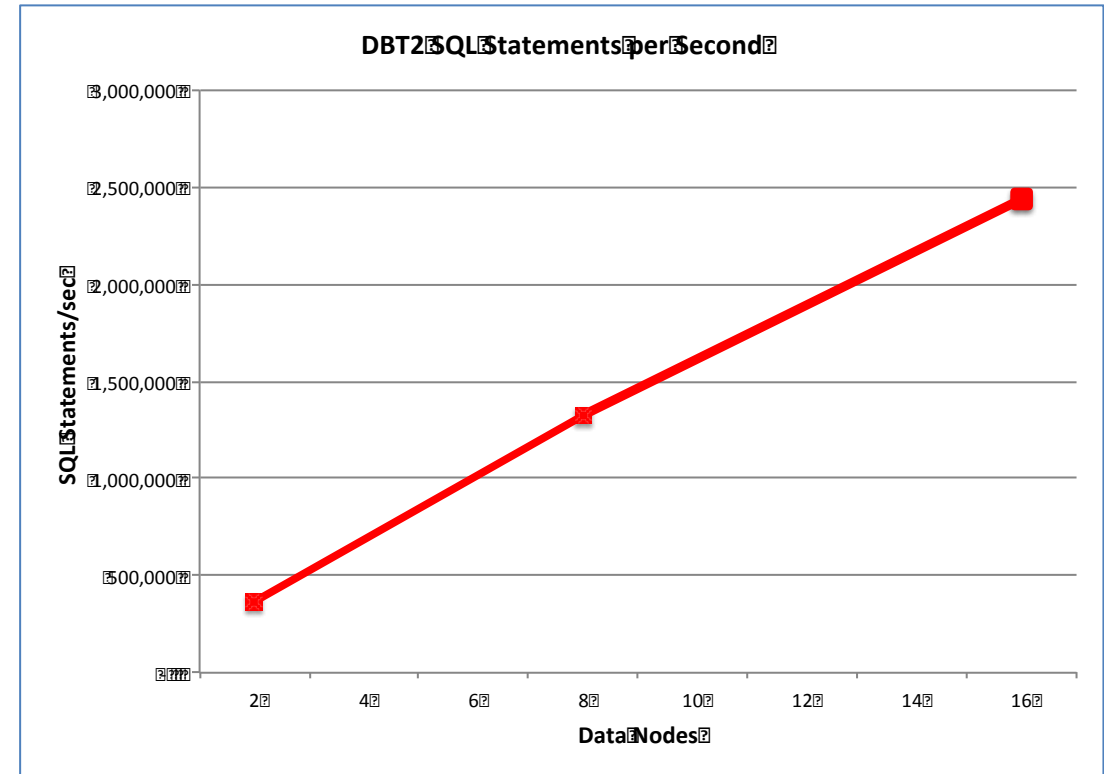
- Memory optimized tables
 - Durable
 - Mix with disk-based tables
- Massively concurrent OLTP
- Distributed Joins for analytics
- Parallel table scans for non-indexed searches
- MySQL Cluster 7.4 FlexAsynch
 - 200M NoSQL Reads/Second



MySQL Cluster 7.4 SQL Performance


2.5M SQL Statements/Second

- Memory optimized tables
 - Durable
 - Mix with disk-based tables
- Massively concurrent OLTP
- Distributed Joins for analytics
- Parallel table scans for non-indexed searches
- MySQL Cluster 7.4 DBT2 BM
 - 2.5M SQL Statements/Second



MySQL Fabric


Scale out with Data Sharding + High Availability



- Scale-out through sharding
 - Read AND Write
 - Standard framework, no more custom solutions
- HA out of the box
 - On top of Replication
 - Automatic failover
 - Automatic routing

Big Data Lifecycle

Better Decisions Using Big Data



Import Data
Apache Sqoop


Apache Sqoop

- Apache TLP, part of Hadoop project
 - Originally developed by Cloudera
- Bulk data import and export
 - Between Hadoop (HDFS) and external data stores
- JDBC Connector architecture
 - Supports plug-ins for specific functionality
- “Fast Path” Connector developed for MySQL

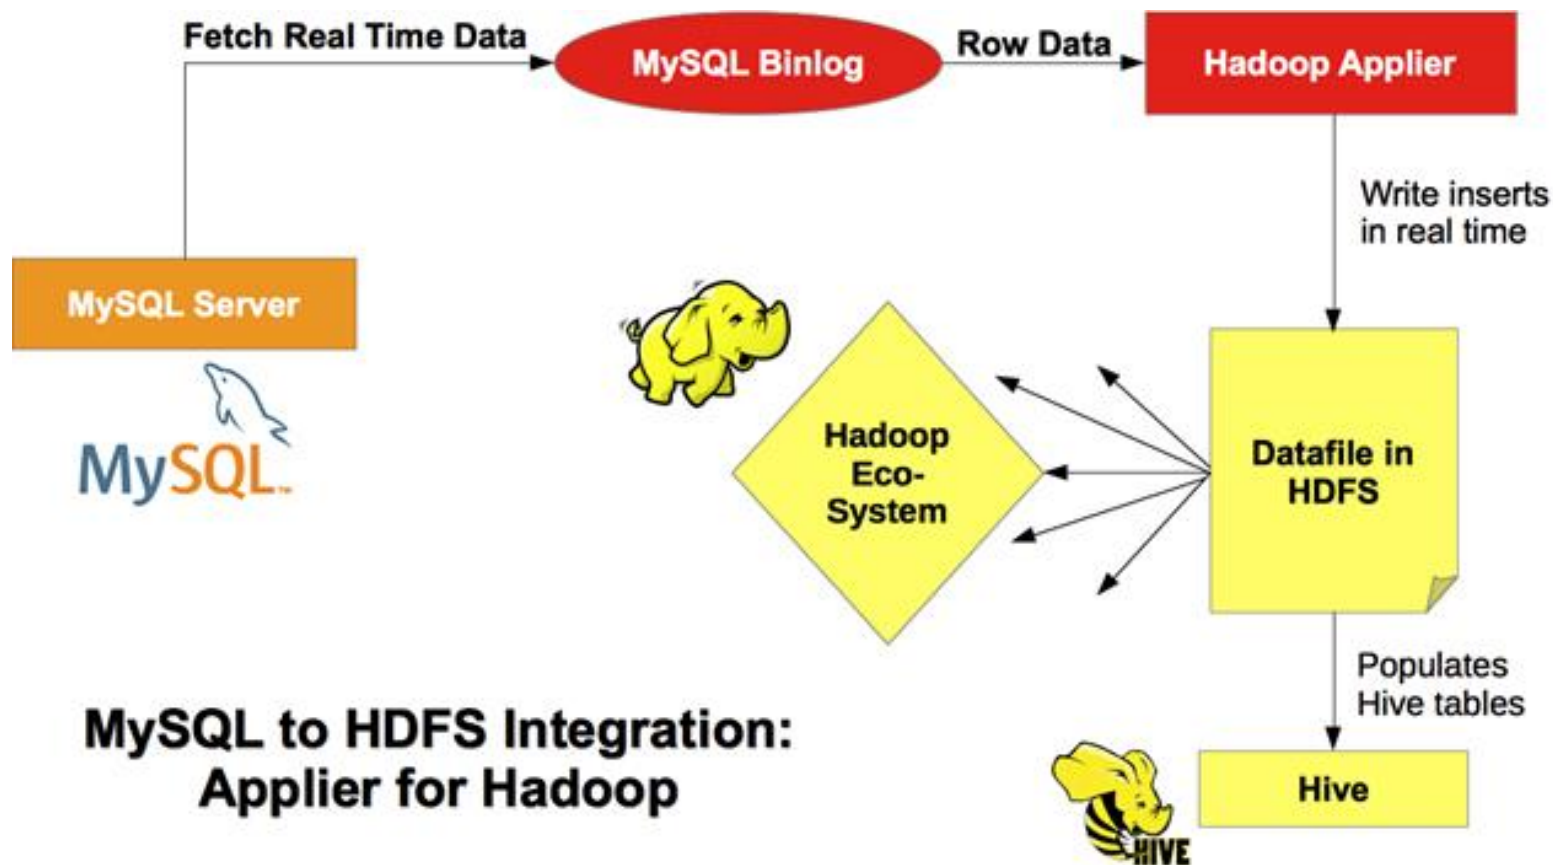


MySQL Applier for Hadoop

- Real-time streaming of events from MySQL to Hadoop
 - Supports move towards “Speed of Thought” analytics
- Connects to the binary log, writes events to HDFS via libhdfs library
- Each database table mapped to a Hive data warehouse directory
- Enables eco-system of Hadoop tools to integrate with MySQL data
- Available for download now: labs.mysql.com




MySQL Applier for Hadoop



Big Data Lifecycle

Better Decisions Using Big Data



Analyze
Export Data
Decide

Analyze Big Data in Hadoop



MySQL Reporting Database for BI



Summary

- Create value from Big Data with MySQL
- MySQL + Hadoop: widely deployed solution (80% of Hadoop project)
- “Best of both worlds”: SQL + NoSQL Access; Schema-less data management
- Scale Out & data sharding with MySQL Fabric
- Tools and expertise to support you



ORACLE®